



DIPARTIMENTO DI **I**NGEGNERIA **E**LETTRICA E **T**ECNOLOGIE DELL'**I**NFORMAZIONE
SCUOLA POLITECNICA E DELLE SCIENZE DI BASE

LAUREA IN INGEGNERIA DELL'AUTOMAZIONE

VLM-Driven Neuro-Symbolic Task and Motion Planning with Ambiguity Resolution

Candidato:
Vincenzo Palomba

Relatori:
Prof. Alberto Finzi

Academic Year 2025/2026

Outline

1. Introduction
2. Deliberative Layer
3. Executive Layer
4. Reactive Layer
5. Setup and Experimental Results
6. Conclusion

Project Objective


A Neuro-Symbolic Task and Motion Planning (TAMP) architecture bridging natural language reasoning with reactive physical execution.

Core Features

- ▶ **Ambiguity Resolution:** Proactively identify and resolve referential ambiguities in human prompts (e.g., "*Which cube?*") via multi-turn visual dialogue prior to physical execution.
- ▶ **Symbolic Verification:** Generate PDDL plans from semantic perception to guarantee causal correctness and eliminate syntactic and causal plan hallucinations given the grounded state.
- ▶ **Dynamic Behavior Tree:** Execute discrete symbolic plans using dynamically generated Behavior Trees, continuous condition monitoring, and auto-recovery fallbacks.
- ▶ **Hybrid Manipulation:** Combine Task-Space RRT for collision-free transit with RL (Proximal Policy Optimization) for compliant, contact-rich grasping.


Scenario: *TaskSorting Arena*

The simulation environment is built on top of the Robosuite library with further customizations.



- ▶ **Robot Hardware:** 7-DoF Franka Emika Panda.
- ▶ **Scene Objects:** Cubes, cans and bins.
- ▶ **Physics Engine:** MuJoCo headless continuous simulation.
- ▶ **Perception:** Dual-view cameras (Top-Down + Frontal).

Neuro-Symbolic Planning architecture



Perception and Ambiguity Resolution

Grounds natural language human instructions in the observed scene and explicitly reasons about task ambiguity through multi-turn visual dialogue, resolving ungrounded assumptions before symbolic deliberation.

Ambiguity Resolution Pipeline

► VLM Visual Scene Grounding (Qwen 2.5-VL):

Parses the camera frame into entity schemas:

- `task_objects`: Manipulable items (e.g., *cans*, *cubes*).
- `target_locations`: Receptacles (e.g., *sorting bin*, *pot*).

► Ambiguity Detection: Evaluates query against detected scene entities:

- Flags referential ambiguity (e.g., *multiple cubes present*).
- Schema: `AmbiguityReasoning(task_ambiguous, question)`.

► Human-in-the-Loop Clarification: Prompts the user with a targeted question, appends response to chat history, and disambiguates the scene.

► Symbols Grounding: Outputs clean symbols $\{\text{Objects}, \text{Locations}\} \subset O_{\text{obj}}, O_{\text{loc}}$ directly to the LLaMA-3 PDDL synthesizer.

Disambiguation Example

Human: *"Put the cube in the container."*

```
[Visual Scene Grounding]
task_objects: [yellow_cube, purple_cube]
target_locations: [sorting_bin, pot]
Ambiguity_Flagged: True
```

Robot: *"Which cube should I move, the yellow or the purple one? And into the pot or the bin?"*

Human: *"The yellow cube into the pot."*

```
[Symbols Grounding]
target: yellow_cube → (type obj)
destination: pot → (type location)
```

VLM Architectures: Dual-Encoder Foundations

Vision-Language Models (VLMs) integrate visual perception with linguistic reasoning by bridging continuous visual feature spaces and discrete language token spaces.




Figure: Standard Dual-Encoder VLM Architecture

Training & Alignment

Trained by aligning modalities using contrastive pretraining or lightweight projection adaptors (MLP) coupled with autoregressive instruction tuning.

Language Encoder (Transformer)

- ▶ Transforms discrete sequences into continuous numerical **embeddings**.
- ▶ Utilizes **self-attention mechanisms** to dynamically weigh token importance.
- ▶ **Decoders** auto-regressively generate the most statistically probable sequence.

Vision Encoder (CNN / ViT)


- ▶ Extracts hierarchical spatial features via Convolutional Neural Networks or Vision Transformers.
- ▶ ViTs segment images into non-overlapping patches, processing them as sequential tokens.

Qwen 2.5-VL Architecture

A state-of-the-art multimodal foundation model utilizing a dynamic-resolution Vision Transformer coupled with a robust language decoder.

Core Components


- ▶ **Qwen2.5 LM Decoder:** Advanced autoregressive language model for reasoning and generation.
- ▶ **Vision Transformer (ViT):** Uses 3D convolutions to efficiently process visual data into sequences of 14×14 patches.
- ▶ **Dynamic Preprocessor:** Adjusts image size dynamically to preserve native aspect ratios and extracts frames optimally for video inputs.



Design Rationale for this Architecture

- ▶ **Dynamic Resolution:** Preserves small object aspect ratios without distortion.
- ▶ **Local Real-Time Inference:** Runs locally via Ollama (1.25 s latency, zero API dependency).
- ▶ **High Spatial Fidelity:** Superior grounding of manipulable items and receptacle boundaries.

Perception: Visual Recognition and Symbolic Grounding



Given the raw overhead RGB camera frame and the natural language instruction, **Qwen 2.5-VL** performs zero-shot visual recognition and symbolic grounding, categorizing required entities into typed PDDL symbols (`obj` vs `location`) while ignoring unreferenced background distractors.

Symbolic Planning: PDDL Formulation

Symbolic Planning uses formal logic to find a sequence of actions that transition a system from an initial state to a target goal state. **PDDL** (Planning Domain Definition Language) is the standard modeling language for this, decoupling the robot's physical capabilities from specific task specifications to guarantee causal correctness before execution.

Formal Planning Tuple

$$\mathcal{P} = (\mathcal{D}, O, s_0, G)$$

► Planning Domain $\mathcal{D} = (\mathcal{L}, \mathcal{A})$ [Static]:

- $\mathcal{L}$: Predicates describing state relations.
- $\mathcal{A}$: Action schemas with preconditions and effects.
- *Fixed and verified offline* (robot capability bounds).

► Planning Problem Instance [Dynamic]:

- O : Grounded universe of objects and locations.
- s_0 : Completely grounded initial state from VLM.
- G : Target goal conditions from user task.
- *Synthesized online* via LLaMA 3 from perception.

Key Principle: Domain physics are fixed; only state and goal are queried online.

Manipulation Domain Definition (domain.pddl)

```
; Grounded State Predicates  $\mathcal{L}$ 
(holding ?obj)
(on-table ?obj)
(on ?obj ?loc)

; Robotic Action Schemas  $\mathcal{A}$ 
(:action pick
  :parameters (?o - obj)
  :precondition (and (on-table ?o)
                     (not (holding ?o)))
  :effect (and (holding ?o)
               (not (on-table ?o))))

(:action place
  :parameters (?o - obj ?loc - location)
  :precondition (holding ?o)
  :effect (and (on ?o ?loc)
               (not (holding ?o))))
```

Neuro-symbolic planning: PDDL Problem Synthesis

While domain physics $\mathcal{D}$ are fixed offline, the planning problem $\mathcal{P} = (\mathcal{D}, O, s_0, G)$ is formulated online. **LLaMA 3** translates visually grounded symbols and task intent into (O, s_0, G) to complete $\mathcal{P}$.

Problem Synthesis Workflow

- ▶ **1. Grounded Few-Shot Prompting:** Uses grounded entities ($O_{\text{obj}}, O_{\text{loc}}$), scene state s_0 , and task G alongside domain $\mathcal{D}$. Enforces strict typing ($O_{\text{obj}} \rightarrow \text{obj}$, $O_{\text{loc}} \rightarrow \text{location}$) and maps task goals to domain predicates.
- ▶ **2. Pydantic Validation, Auto-Repair & Compilation:** Validates LLM output against pydantic Problem schema, auto-correct undeclared symbols, and retries on predicate violations. `Problem.to_pddl()` deterministically serializes the verified instance into standard PDDL syntax.


Key Guarantee: Structured JSON + Pydantic validation eliminates syntactic and schema hallucinations in the problem instance given the grounded state.

Generated Planning Problem (problem.pddl)

```
; Online problem instance  $\mathcal{P} = (\mathcal{D}, O, s_0, G)$ 
(define (problem manipulation-task)
  (:domain manipulation)
  (:objects
    red_can yellow_cube blue_can - obj
    sorting_bin pot - location
  )
  (:init
    (on-table red_can)
    (on-table yellow_cube)
    (on-table blue_can)
  )
  (:goal (and
    (on red_can sorting_bin)
    (on yellow_cube pot)
    (on blue_can sorting_bin)
  ))
)
```

Plan Verification: Fast Downward Solver

Fast Downward acts as the symbolic deliberation engine, using initial states s_0 , final goal G and Domain specification $\mathcal{D}$ to obtain a verified action sequences via **Heuristic Forward State-Space Search**.



Deliberation Guarantee: Computes verified causal action plans in < 0.08 s, eliminating syntactic and causal plan hallucinations given the grounded state.

Forward Search & Heuristics

- **Search Paradigm:** Progresses forward in state space $\mathcal{S}$:

$$s_0 \xrightarrow{a_1} s_1 \rightarrow \dots \rightarrow s_g \models G.$$

- **Direct Grounding:** Grounded s_0 from VLM makes forward state progression directly evaluable.
- **State Transitions:** Applies applicable action a ($\text{pre}(a) \subseteq s$):
 $s' = (s \setminus \text{del}(a)) \cup \text{add}(a)$.
- **Heuristic Engine (LAMA / h_{FF}):**

- **Lazy GBFS:** Greedy best-first search for high-speed plan discovery.
- **Dual Heuristics ($h_{FF} + h_{LM}$):** Delete-relaxation ($\text{del}(a) = \emptyset$) + causal Landmarks.
- **Preferred Operators:** Prunes search branching via helpful actions.

PDDL Problem Generation experiments

To bridge perceptual symbols and symbolic planning, the deliberation layer couples **Pydantic data models** with programmatic domain validation and a reflection loop to eliminate syntactic and schema hallucinations given the grounded state.

Structured Pipeline & Pydantic Schemas

- ▶ **Synthesizer Class** (`AmbresStructured`):
`generate_problem_json(...)` queries LLaMA 3 with $(s_0, G, O_{\text{obj}}, O_{\text{loc}}, \mathcal{D})$ under few-shot JSON constraints.
- ▶ **Data Models** (`src.llm2pddl.problem`):
 - ▶ **Problem**: Declares objects, init, and goal.
 - ▶ **PDDLObject**: Type enforcement (obj, location).
 - ▶ **Predicate**: Structured facts with arity validation.
- ▶ **Compilation** (`Problem.to_pddl()`):
Deterministically convert the JSON payload schema into standard PDDL formatted text.

Closed-Loop Validation & Auto-Repair

- ▶ **Validation Function**:
`generate_problem_with_retry(..., \mathcal{D})` verifies predicates and types against $\mathcal{D}$ via regex extraction.
- ▶ **Handled Failure Modes**:
 - ▶ *Predicate Hallucination*: e.g., $(\text{in } ?x \text{ ?bin}) \notin \mathcal{D}$.
 - ▶ *Undeclared Symbols*: Missing object in objects.
 - ▶ *Type Inversion*: Receptacle labeled as obj.
- ▶ **Auto-Healing & Reflection**:
 - ▶ Inferred auto-declaration for missing receptacles.
 - ▶ Re-prompts LLaMA 3 with exact error traces (max 5 retries).

PDDL Problem Generation experiments

Empirical validation across **28 simulation trials** demonstrates that closed-loop PDDL synthesis achieves **100.0%** feasibility and reliably shields the executive layer.

Empirical Benchmark

Evaluation Metric	Result
Total Logged Trials	28
1st-Pass Solver Feasibility	26 / 28 (92.9%)
Success after Feedback Loop	28 / 28 (100.0%)
Fast Downward Parser Errors	0% (0 / 28)
Avg. Symbolic Solve Time	0.086 s
Max Sequential Objects	4 (8 plan steps)

- ▶ **92.9% First-Pass Feasibility:** Raw LLaMA 3 output before verification.
- ▶ **100% Closed-Loop Success:** Feedback reflection repairs 100% of schema errors (2/2) on Retry 1.


Impact on the Executive Layer

- ▶ **Causal Feasibility Verification:**
Fast Downward verifies task feasibility *prior* to actuation; unreachable goals fail safely offline.
- ▶ **Plan-to-BT Compiler:**
`build_bt_from_pddl_plan(plan, adapter)` compiles verified plans $\pi = \langle a_1, \dots, a_N \rangle$ into reactive Behavior Trees.
- ▶ **Executive Shielding:**
The Behavior Tree is shielded from malformed symbols, broken preconditions, or causally invalid action sequences given the grounded state.

Bridge to Execution: Formal symbolic deliberation ensures causal and syntactic plan consistency given the grounded state before handoff to Behavior Tree execution.



Executive Layer: Behavior Trees

- ▶ A **Behavior Tree** is a directed tree with nodes classified as root, control flow, or execution (tasks).
- ▶ Execution starts from the root, which sends *ticks* to its children at a set frequency.
- ▶ A node returns one of three statuses: **Running**, **Success**, or **Failure**.



Control Flow Nodes


- ▶ **Selector:** A fallback node returns **Success** if at least one child do not fail. If every child fail return **Failure**. Otherwise **Running**.
- ▶ **Sequence:** Executes children in order, returning success if everyone succeeds. Returns **Failure** or **Running** immediately if a child fails or is still running.



How to build Behavior Trees: `py_trees` library

`py_trees` is a Python library used to construct Behavior Trees efficiently. It relies primarily on two base classes: **Behaviours** and **Composites**.

- ▶ **Behaviour (Leaf):** The smallest element in a tree. Behaviours typically represent a **condition** (e.g., *Am I hungry?*) or an **action** (e.g., *Buy chocolate cookies*).
- ▶ **Composites (Control Flow):** Responsible for directing the path traced through the tree during a tick.
 - ▶ **Selectors:** Fallback decision-makers.
 - ▶ **Sequences:** Order-dependent factories.
 - ▶ **Parallels:** Concurrent execution factories.



BT Builder: Parsing and Sequence Compilation

1. Parsing PDDL Plan

(`parse_pddl_action`) Cleans raw planner text into structured Python tokens:

- ▶ **Syntax Stripping:** Removes step indices, outer parentheses, and comment lines (`; cost = 2`).
- ▶ **Token Extraction:** Normalizes action name and collects args:
`"0: (place yellow_cube pot)" → ("place", ["yellow_cube", "pot"])`.

2. Plan to Sequence

(`pddl_plan_to_sequence`) Compiles actions into a closed-loop executable Sequence:

- ▶ **Skill Dispatch:** Maps actions via `SKILL_REGISTRY` to motion primitives (`PickUp`, `PlaceInBin`).
- ▶ **Action Post-Conditions (PDDLCondition):** Wraps each action into a sub-sequence [Action → PostCondition]:
 - `pick → PickSeq: PickUp → Holding(o)`
 - `place → PlaceSeq: PlaceInBin → On(o, loc)`
- ▶ **Fail Verification:** If grasp or placement fails, post-condition returns `FAILURE` immediately, halting blind execution before cascading errors occur.

Compilation Pipeline Example

1. Planner Output (Fast Downward):

```
; plan generated by fast downward
0: (pick yellow_cube)
1: (place yellow_cube pot)
```

↓ *parsed & compiled into BT*

2. Executable PyTrees Sequence:

```
[ - ] pddl_plan_sequence (memory=True)
| -- [ -> ] PickSeq(yellow_cube)
|   | -- PickUp(object='yellow_cube')
|   |   -- [>] Holding(yellow_cube)      <-- Post-Cond
|   | -- [ -> ] PlaceSeq(yellow_cube->pot)
|   |   | -- PlaceInBin(yellow_cube->pot)
|   |   |   -- [>] On(yellow_cube->pot)      <-- Post-Cond
```

Design Rationale

Decouples **symbolic planning** from **continuous control**, while enforcing **closed-loop post-condition checks** after every single action.

BT Builder: Goal Checking and Tree Generation

3. Goal-Guarded Wrapper

(`wrap_with_goal_check`) Augments the plan sequence with closed-loop reactive guards:

- ▶ **Pre-Check** (`GoalCheck_Pre`): Verifies if goal is already met before acting → succeeds with 0 redundant actions.
- ▶ **Action Post-Conditions**: Intercepts failures locally at each action step (`Holding`, `On`) before downstream steps run.
- ▶ **Post-Check** (`GoalCheck_Post`): Confirms overall physical goal completion $\bigwedge g_i$ after sequence finishes.
- ▶ **Retry Loop** (`Retry`): If any action post-condition or post-check fails, automatically retries the plan up to N times.

4. Main Entrypoint

(`build_bt_from_pddl_plan`) Single builder API providing two execution modes:

- ▶ `wrap_goal_check=False`: Sequence with action post-conditions.
- ▶ `wrap_goal_check=True`: Self-healing tree with retry loop.

Goal-Guarded Tree Architecture

```
[?] goal_guarded_root (Selector)
|-- [>] GoalCheck_Pre (already met?)
'-- [R] retry_until_goal (max=10)
    '-- [->] attempt (Sequence)
        |-- [->] pddl_plan_sequence
        |   |-- [->] PickSeq(yellow_cube)
        |   |   |-- PickUp(yellow_cube)
        |   |   '-- [>] Holding(yellow_cube)
        |   '-- [->] PlaceSeq(cube->pot)
        |       |-- PlaceInBin(cube->pot)
        |       '-- [>] On(cube->pot)
        '-- [>] GoalCheck_Post (verified?)
```

Failure Recovery Example (Cube Slip)

1. Robot picks up cube, but gripper slips mid-air.
2. Action post-condition `Holding` fails immediately.
3. `PickSeq` aborts → prevents blind place attempt.
4. Retry catches failure → re-executes plan → task completed!

BT Nodes: Skills & Conditions

The executive layer instantiates action skills and condition monitors as modular `py_trees` nodes, bridging discrete PDDL plans to continuous 20 Hz robot control.

Action Skills (`skills.py`)

- ▶ **Base Architecture (Skill):**
Subclasses `Behaviour`. Non-blocking
20 Hz tick: returns `RUNNING`, `SUCCESS` on `is_done()`, or `FAILURE` on timeout. Emits 7-DOF OSC commands $a \in \mathbb{R}^7$.
- ▶ **Pick Primitive (MotionPlanningPickUp):**
4-stage FSM: `follow_path` (RRT to hover) $\rightarrow$ `descend` (linear approach) $\rightarrow$ `grasp` (finger closure) $\rightarrow$ `lift` (clearance).
▶ *Hybrid PPO Variant:* Replaces descent/grasp with learned RL policy.
- ▶ **Place Primitive (MotionPlanningPlaceInBin):**
2-stage FSM: `follow_path` (RRT transit with locked grip) $\rightarrow$ `release` (open gripper over container).

Condition Nodes (`conditions.py`)

- ▶ **Post Conditions Checking(conditions.py):**
Evaluates actions `add(s)` effects at zero latency:
 - `on(?x, ?y)  $\leftrightarrow$  env.is_in_bin(obj, target)`
 - `holding(?x)  $\leftrightarrow$  env.is_grasped(obj)`
 - `on-table(?x)  $\leftrightarrow$   $\neg$ grasped  $\wedge$   $\neg$ in_bin`
- ▶ **Goal Monitor (PDDLGoalCheck):**
Parses conjunctive goals $\bigwedge g_i$ from `problem.pddl`. Evaluates pre/post conditions to skip met goals or trigger retry.

Execution & Runtime Tree

The runtime engine executes the compiled Behavior Tree via a 20 Hz tick loop, continuously verifying action post-conditions, dispatching motion skills, and recovering from disturbances.

Runtime Execution Pipeline

- ▶ **20 Hz Tick Loop (py-trees):**
Synchronous active-leaf evaluation; steps MuJoCo with 7-DOF OSC commands $a \in \mathbb{R}^7$.
- ▶ **Closed-Loop Reactive Flow:**
 - *Pre-Check:* GoalCheck.Pre skips met goals (0 redundant steps).
 - *Skill Dispatch:* Steps TaskSpaceRRT waypoints or PPO policy.
 - *Post-Check:* Holding / On verify physical effect immediately.
 - *Recovery:* Action failure triggers retry_until_goal loop.
 - *Post-Check:* GoalCheck.Post verify if all the goal predicates are true.
- ▶ **Real-Time Telemetry:** Dual-camera HUD overlays active leaf, skill status, and ticks.


Goal-Guarded Tree Architecture

```
[?] goal_guarded_root (Selector)
|-- [>] GoalCheck_Pre (already met?)
'-- [R] retry_until_goal (max=10)
    '-- [->] attempt (Sequence)
        |-- [->] pddl_plan_sequence (Memory)
        |   |-- [->] PickSeq(yellow_cube)
        |   |   |-- PickUp(yellow_cube)
        |   |   |-- [>] Holding(yellow_cube)
        |   |   |-- [->] PlaceSeq(cube->pot)
        |   |   |-- PlaceInBin(cube->pot)
        |   |   |-- [->] On(cube->pot)
        |   '-- [->] GoalCheck_Post (verified?)
```

- ▶ **Selector [?]:** Skips plan if met; falls back to retry loop.
- ▶ **Sequence [->]:** Action $\rightarrow$ Post-check; halts on slip/drop.
- ▶ **Retry [R]:** Autonomous recovery loop (≤ 10 attempts).

Reactive Robustness: Closed-loop verification of positive post-conditions halts flawed execution immediately, enabling autonomous recovery without cascading errors.

Main Scenario: Generated Behavior Tree



1. Goal Conditions Check

- **GoalCheck_Pre**: skips action if goal met (0 latency).
- **GoalCheck_Post**: verifies overall task formula $\mathcal{G} = \bigwedge_i on(x_i, loc_i)$.

2. Action Sequence

- **Pick**: pick action and effects conditions $Holding(x)$ ($\Delta z > 2.5$ cm) checking.
- **Place**: place action and effects conditions $On(x \rightarrow loc)$ (containment) Checking.
- Halts execution immediately on FAILURE conditions (slip/miss).

3. Reactive Recovery

- **retry_until_goal** ≤ 10 attempts.
- Intercepts disturbance $\rightarrow$ re-attempts plan rather than blindly continuing.


Reactive Layer: Motion and Path Planning

Reactive Layer Roles

- ▶ **Symbolic-to-Continuous Grounding:** Translates discrete BT skills (pick, place) into continuous robot motion.
- ▶ **Perception:** Accounts for physical geometry, obstacles, and contact dynamics ignored by PDDL.
- ▶ **Motor Control:** Operates at 20 Hz with real-time state feedback from MuJoCo Simulator Environment.

Hierarchical Motion Decomposition

- ▶ **1. Path Planning (Task Space RRT):** Solves pure geometric obstacle avoidance in free space ($\mathbb{R}^3$).
- ▶ **2. Trajectory Tracking:** Closed-loop proportional tracking of Cartesian waypoints $[x_1, \dots, x_n]$ generated by the RRT.
- ▶ **3. Manipulation (PPO / Heuristic):** Governs fine contact dynamics, compliant grasping, and lifting.
- ▶ **4. Low Level Controller (OSC):** Takes as input the cartesian velocities from the Tracking and Manipulation modules and output the joint torques τ_i needed to control the robot.



Path Planning: RRT

Rapidly Random Trees (RRT) is a Probabilistic Path Planning method used to efficiently construct a path in continuous spaces.

Algorithm Goal Oriented RRT Algorithm

```
1:  $G.\text{init}(x_{\text{start}})$ 
2: for  $i = 1 \dots K$  do
3:    $p = \text{UniformRandom}(0, 1)$ 
4:   if  $p \leq p_{\text{goal}}$  then
5:      $x_{\text{rand}} \leftarrow x_{\text{goal}}$ 
6:   else
7:      $x_{\text{rand}} \leftarrow \text{RandPose}()$ 
8:   end if
9:    $x_{\text{nearest}} \leftarrow \text{NearestVertex}(x_{\text{rand}}, G)$ 
10:   $x_{\text{new}} \leftarrow \text{NewConf}(x_{\text{nearest}}, x_{\text{rand}}, \Delta x)$ 
11:  if  $\text{CollisionFree}(x_{\text{nearest}}, x_{\text{new}})$  then
12:     $G.\text{add\_vertex}(x_{\text{new}})$ 
13:     $G.\text{add\_edge}(x_{\text{nearest}}, x_{\text{new}})$ 
14:    if  $x_{\text{new}} \in X_{\text{goal}}$  then return Path Found
15:  end if
16: end if
17: end for
18: return  $G$ 
```

The goal oriented version samples the goal point with a given probability p while a random point with probability $1 - p$.

Trajectory Tracking

- ▶ **Cartesian Setpoints:** The waypoints generated by the task-space RRT are Cartesian poses used as setpoints for a **Proportional (P) Controller**.
- ▶ **Tracking Error:** At each time step, the position error is defined as:

$$\mathbf{e} = \mathbf{p}_{\text{target}} - \mathbf{p}_{\text{EEF}}$$

- ▶ **Control Law:** The controller outputs velocity commands bounded by a maximum speed:

$$\mathbf{v}_{\text{cmd}} = \text{clip}(K_p \mathbf{e}, -v_{\text{max}}, v_{\text{max}})$$


which are sent to the underlying operational space controller.

- ▶ **Application:** Used specifically during the **approach phase** in the pick action and during the **place action**. The place skill itself is completely based on the RRT planner and the trajectory tracker, while the pick skill is composed by more modules and need further explanation.

Manipulation: Heuristic Approach

Consists of a **Staged FSM**, that executes the following phases:

1. **RRT Planning:** Traverse collision-free waypoints to hover pose ($== \text{end_pose}$).
2. **Descent Phase:** Descend vertically until reaching grasp height ($z \leq z_{\text{tol}}$).
3. **Grasp Phase:** Actuate gripper firmly for set duration ($t > t_{\text{grasp}}$).
4. **Lift Phase:** Elevate object to safe height to verify pick. If it fails, it will retry from the Descent phase.



Manipulation: Proximal Policy Optimization

Proximal Policy Optimization is an on-policy actor-critic Reinforcement Learning method. PPO updates policy parameters θ_{k+1} by maximizing the **Clipped Surrogate Objective**:

$$L^{\text{CLIP}}(\theta) = \hat{\mathbb{E}}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right]$$

where the probability ratio is defined as:

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_k}(a_t | s_t)}$$

- ▶ **Clipped Objective:** The hyperparameter $\epsilon \approx 0.2$ penalizes moving $r_t(\theta)$ outside $[1 - \epsilon, 1 + \epsilon]$, preventing destructively large policy updates.
- ▶ **Actor-Critic Architecture:** Jointly optimizes a policy network $\pi_{\theta}(a|s)$ and a state-value baseline $V_{\phi}(s)$ using Generalized Advantage Estimation (GAE).

Manipulation: Observation Space

TaskSpaceRRT provides coarse global navigation to a pre-grasp hover pose (~ 18 cm above the object). The **PPO policy** then takes over using a minimal, translation-invariant **11D observation vector**:

$$\mathbf{o}_t = \left[\underbrace{\Delta \mathbf{p}}_{\mathbb{R}^3}, \underbrace{\mathbf{q}_{\text{grip}}}_{\mathbb{R}^2}, \underbrace{\dot{\mathbf{q}}_{\text{grip}}}_{\mathbb{R}^2}, \underbrace{z_{\text{rel}}^{\text{EEF}}}_{\mathbb{R}^1}, \underbrace{\Delta z_{\text{obj}}}_{\mathbb{R}^1}, \underbrace{c_{\text{grasp}}}_{\{0,1\}}, \underbrace{m_{\text{geom}}}_{\{0,1\}} \right]^T \in \mathbb{R}^{11}$$

Kinematics & Spatial Relations

- ▶ $\Delta \mathbf{p} = \mathbf{p}_{\text{obj}} - \mathbf{p}_{\text{EEF}}$: 3D relative position vector (translation invariance across table).
- ▶ $z_{\text{rel}}^{\text{EEF}} = p_{\text{EEF},z} - z_{\text{table}}$: Table clearance to avoid surface collisions.
- ▶ $\Delta z_{\text{obj}} = p_{\text{obj},z} - p_{\text{obj},z}^{(0)}$: Vertical displacement measuring lift progress.

Gripper and Auxiliary States

- ▶ $\mathbf{q}_{\text{grip}}, \dot{\mathbf{q}}_{\text{grip}}$: Left/right finger joint positions and velocities (contact stall detection).
- ▶ $c_{\text{grasp}} \in \{0, 1\}$: Bilateral force-closure flag (Robosuite contact normal check).
- ▶ $m_{\text{geom}} \in \{0, 1\}$: Shape token (0: cube, 1: can) enabling multi-geometry conditioning.

Manipulation: Reward Function

The agent is trained with a multi-stage continuous-discrete reward guiding the transition from approach $\rightarrow$ grasp $\rightarrow$ lift:

$$R_t = r_{\text{reach}} + r_{\text{grip}} + r_{\text{grasp}} + r_{\text{lift}} - c_{\text{step}} + r_{\text{terminal}}$$

Continuous Shaping Terms

- ▶ **Smooth Reaching** ($d = \|\Delta \mathbf{p}\|_2$):
 $r_{\text{reach}} = 1.0 - \tanh(10.0 \cdot d)$
- ▶ **Grasp Maintenance:**
 $r_{\text{grasp}} = 2.0 \cdot \mathbb{I}(\text{grasped})$
- ▶ **Conditioned Vertical Lift:**
 $r_{\text{lift}} = 20.0 \cdot \max(0, \Delta z_{\text{obj}}) \cdot \mathbb{I}(\text{grasped})$
(Avoids lifting with an empty gripper)
- ▶ **Alive Cost:** $c_{\text{step}} = 0.01$ (incentivizes speed)

Action Shaping & Terminals




- ▶ **Gripper Timing Incentive** r_{grip} :

$$r_{\text{grip}} = \begin{cases} +0.5 & \text{if } d < 4 \text{ cm} \wedge a_g > 0 \\ -0.25 & \text{if } d < 4 \text{ cm} \wedge a_g \leq 0 \\ -0.2 & \text{if } d \geq 4 \text{ cm} \wedge a_g > 0.2 \\ 0 & \text{otherwise} \end{cases}$$

Prevents premature closure before alignment.

- ▶ **Terminal Sparse Rewards:**
 $+15.0$ if $\Delta z_{\text{obj}} \geq 6 \text{ cm} \wedge \text{grasped}$
(Success)
 -2.0 if $p_{\text{obj},z} < z_{\text{table}} - 4 \text{ cm}$ (Drop)

Manipulation: PPO Training & Evaluation




Training Return

- **Stochastic Action Noise:** Persistent Gaussian exploration noise (with standard deviation $\sigma \approx 1.0$) perturbs micro-contacts during training (~ 50 avg).
- **Peak Episodes (≈ 200 – 230):** When random noise aligns perfectly, trajectories achieve full force closure and lift with terminal bonus (+15).

Evaluation and Task Completion

- **Deterministic Evaluation:** Setting $\sigma = 0$ and evaluating via the distribution mean μ effectively removes action jitter, revealing the true learned policy return of **199.4** (near theoretical max ~ 200).
- **Success Rate Ascent:** Climbs rapidly from **0% $\rightarrow$ 45% $\rightarrow$ 95%** robust task completion.

Manipulation: Behavioral Cloning Warm-Start




Synthetic Demonstrations

- ▶ Generated via a closed-loop 3-stage state-machine using simulation ground-truth poses.
- ▶ 100 demos ($\sim 3,000$ transitions) pre-trained for 150 epochs via supervised MSE loss.

Pure RL vs BC-RL

- ▶ **Pure RL (120k steps):** Gets stuck in the "*reaching local minimum*" (≈ 29 reward), never discovering the precise descent-grasp-lift sequence.
- ▶ **BC Warm-Start (25k steps):** Initializes policy in the basin of attraction, achieving immediate task competence in $5\times$ fewer steps.

Manipulation: Heuristic vs PPO Comparison



Manipulation Performance

- **Generalization:** PPO achieves **95%** on **Yellow Cube** (consistent with training) and **90%** on **Purple Cube** (zero-shot transfer), averaging **92.5%**.
- **Micro-Manipulation Latency:** PPO median is **1.10 s** (22 steps) vs. Heuristic **1.72 s** (34.3 steps) (PPO mean is 33.5 due to rare retry outliers).
- **Eliminates Dwell Timers:** Heuristic relies on fixed post-grasp dwell timers; PPO senses force closure and lifts dynamically upon contact.

Trade-Offs

- **Dynamic Impedance:** PPO modulates continuous velocities based on real-time force closure rather than rigid fixed gains ($K_p = 10$).
- **Privileged vs Reactive:** Heuristic requires exact object coordinates; PPO relies only on local $\Delta \mathbf{p}$ and joint kinematics.
- **Safety Guard:** Hybrid pick node falls back to heuristic if step budget (≤ 90) expires.

Experimental Setup: Simulation Environment



Robosuite TaskSorting Arena

- ▶ **Robot:** 7-DoF Franka Emika Panda with parallel-jaw gripper (20 Hz OSC control).
- ▶ **Physics:** MuJoCo headless physics engine with compliant contact dynamics.
- ▶ **Dual Viewpoints:** Orthographic *top-down* (VLM grounding) + perspective *frontview* (monitoring).
- ▶ **Objects:** 5 pickable items (yellow/purple cubes, red/blue/green cans).
- ▶ **Containers:** Sorting bin (30×30 cm) + handled pot (collision obstacles).

End-to-End TAMP Pipeline Execution

Task Mission: *"Put the yellow cube in the pot, then put the red can in the sorting bin, then put the blue can in the sorting bin."*

1. Perception & Plan

- ▶ **Qwen 2.5-VL:** Scene Detection.
- ▶ **Llama 3:** PDDL synthesis.
- ▶ **Fast Downward:** < 0.1 s plan.

2. BT Compilation

- ▶ Compiled py_trees reactive tree.
- ▶ Pre-conditions & fallback nodes.
- ▶ 20 Hz continuous monitoring.

3. Motion & Control

- ▶ **TaskSpaceRRT:** Paths (~ 30 ms).
- ▶ **Hybrid PPO:** Grasp/lift (≈ 1.7 s).
- ▶ **Placement:** Waypoint tracking.



▶ Heuristic Trial (simulation_trial_heuristic.gif)



▶ PPO Policy Trial (simulation_trial_ppo.gif)

▶ GitHub Repo: [vincip99/.../tree/master](https://github.com/vincip99/tamp)

Experimental Results: Main Scenario Performance

Pipeline Layer	Module / Algorithm	Latency / Steps
Visual Grounding	Qwen 2.5-VL (3B)	1.25 s
PDDL Formulation	LLaMA 3 (8B)	0.65 s
Symbolic Planning	Fast Downward (A^*)	0.08 s (6 causal actions)
Motion Planning	TaskSpaceRRT	0.03 s / transit leg (6 legs)
Manipulation (Pick)	Macro Transit + PPO Grasp*	79–132 steps / pick
Reactive Executive	Behavior Tree (py_trees)	445 ticks (22.2 s)
End-to-End System	Complete Architecture	24.2 s Total (100% Success)

Planning and Execution

- ▶ **Causal Plan:** $\langle \text{pick}(\text{cube}) \rightarrow \text{pot}, \text{pick}(\text{red}) \rightarrow \text{bin}, \text{pick}(\text{blue}) \rightarrow \text{bin} \rangle$.
- ▶ **Pick Latency (*):** End-to-end pick combines macro RRT approach transit (~ 55 – 100 steps) and micro PPO grasp (20 – 25 steps, Slide 29).
- ▶ **Total Duration:** Completed in 445 ticks (22.2 s) without timeouts.

Robustness

- ▶ **Obstacle Clearance:** 0 collisions with bin rims or pot handles across all 6 transit trajectories.
- ▶ **Grasp Stability:** Zero object slip or drop events during transit; firm force closure maintained throughout.

Experimental Results: 28-Task Benchmark Performance

Deliberation and execution benchmark across **28 manipulation tasks**:

Benchmark Complexity Tier	Trials	Plan Length	1st-Pass Feas.	Post-Refl. Feas.	Task Execution Success
Tier 1: Single-Object Relocation	5	2 steps	100.0% (5/5)	100.0% (5/5)	100.0% (5/5)
Tier 2: Two-Object Sequential	6	4 steps	83.3% (5/6)	100.0% (6/6)	100.0% (6/6)
Tier 3: Three-Object Chained	7	6 steps	100.0% (7/7)	100.0% (7/7)	100.0% (7/7)
Tier 4: Four-Object High-Complexity	3	8 steps	66.7% (2/3)	100.0% (3/3)	100.0% (3/3)
Tier 5: Ambiguity & Lexical Synonyms	7	2-6 steps	100.0% (7/7)	100.0% (7/7)	100.0% (7/7)
Total Benchmark Dataset	28	Up to 8 steps	92.9% (26/28)	100.0% (28/28)	100.0% (28/28)

End-to-end task completion rate across the 28 empirical tabletop manipulation benchmark scenarios.

LLM Failure Modes & Auto-Repair

- **Schema Validation (7.1%):** 2 syntax/formatting errors (Trials 11 & 21) caught by Pydantic validator.
- **Syntactic Consistency:** Domain predicates and type signatures strictly preserved across all tiers.
- **Auto-Repair Gain:** Lifted feasibility from **92.9% → 100.0%** (2/2 errors resolved on Retry 1).

Manipulation: PPO Failures & BT Fallback

- **PPO Competence: 92.5%** pick success on cubes (**95%** Yellow trained, **90%** Purple zero-shot).
- **PPO Failure Modes:** Out-of-distribution tall cylinder geometry and micro-contact roll cause step budget timeout (≥ 90 steps).
- **Fallback Guarantee:** Behavior Tree fallback triggers deterministic FSM upon timeout → **100%** physical execution success on all dispatched plans (6/6 in Tier 3, 27/27 overall).

Conclusions & Future Perspectives

A hierarchical 3-layer architecture bridging multimodal foundation models with formal symbolic deliberation and closed-loop reactive manipulation.

Core Contributions

- ▶ **Neuro-Symbolic Deliberation:** Autonomous extraction of (s_0, G) from RGB frames and language, fortified by interactive disambiguation and auto-repair.
- ▶ **Formal Fast Downward Gatekeeper:** Guarantees causal plan validity in < 0.08 s before physical dispatch.
- ▶ **Reactive Executive Layer:** Behavior Trees compile action sequences into tick-based subtrees with continuous pre/post-condition monitoring.
- ▶ **Robust Hybrid Control:** TaskSpaceRRT ensures collision-free transit, while BC-bootstrapped PPO achieves reliable contact-rich grasping.

Future Developments

- ▶ **Domain & Problem Co-Synthesis:** Moving beyond fixed domain templates toward online predicate invention and automated action schema discovery.
- ▶ **Skill Registry Augmentation:** Expanding the primitive library to articulated objects and different actions.
- ▶ **Unified RL/IL for TAMP:** End-to-end hierarchical policy learning bridging task logic with real-time visuomotor control.
- ▶ **Sim-to-Real Transfer:** Hardware validation on physical 7-DOF Franka Emika Panda setups.

Key References



E. Chisari, J. O. von Hartz, F. Despinoy, and A. Valada.
Robotic Task Ambiguity Resolution via Natural Language Interaction.
arXiv preprint arXiv:2504.17748 [cs.RO], 2025.



B. Liu, Y. Jiang, X. Zhang, Q. Liu, S. Zhang, and P. Stone.
LLM+P: Empowering Large Language Models with Optimal Planning Proficiency.
arXiv preprint arXiv:2304.11477 [cs.AI], 2023.



N. Wake, A. Kanehira, J. Takamatsu, K. Sasabuchi, and K. Ikeuchi.
VLM-driven Behavior Tree for Context-aware Task Planning.
arXiv preprint arXiv:2501.03968 [cs.RO], 2025.



J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov.
Proximal Policy Optimization Algorithms.
arXiv preprint arXiv:1707.06347 [cs.LG], 2017.